

## Homework 5B

Faculty: Shavit

TA: Coulombe, Saribekyan

## Programming

This assignment is all about high-performance concurrent hash tables. You will build a set of different hash tables and extend one of your choosing in a couple non-trivial ways. Each design will support `ADD()`, `REMOVE()` and `CONTAINS()` methods (as described on page 299 in the text) and will be resizable, based on a *fullness* criteria. That is, you will specify some measure of the occupancy of the table that will trigger a `RESIZE()` operation, which will then double the space and copy the data. There are many such criteria (*eg.* max total number of items in the table, max total items in any given bucket etc.) and for this assignment you will decide which heuristic to use and provide the reasoning and experimental data (of your own design) to support your decision.

The world is awash in variations on hash tables - in this assignment we will focus on two canonical types: **Closed-Address** hash tables have linked lists located at each bucket (or hash index), which support the `ADD()`, `REMOVE()` and `CONTAINS()` methods. **Open-Address** hash tables have individual entries at each hash index, but resolve collisions by inserting items at nearby locations. Please see Chapter 13 in the text for more detail.

All of the hash table designs use locks and while you're permitted to use your lock designs from *pset4*, you are encouraged to make use of the `REentrantLock` provided in the Java concurrency package which provides reliably good performance and nice features (*eg.* `ReadWriteLock` and `TryLock` etc.). In general, we will use the `StripedHashSet` paradigm (Figure 13.6 in the text), allocating a lock bank of size  $2^{\lceil \log_2 n \rceil}$  (the next larger power of two) for  $n$  threads. This allows us to map a hash index to a particular lock with a bitwise AND operation, which is cheap and convenient. Further, in practice it is generally accepted that if the number of locks is approximately equal to the number of threads (with a uniform distribution of accesses to each) you will largely mitigate lock contention. Feel free to demonstrate to yourself that this is true in our case, but it is not necessary for this assignment; you are free to assume it. In addition, we'll make the simplifying assumption that all hash tables have  $2^k$  buckets for some  $k$  - after all, we're not savages.

**Assignment 1. Hash Tables** We will build each of the following types of hash table,  $H$ :

- a. **Lock-based Closed-Address (LOCKING)**: This is a standard hash table where `ADD()` and `REMOVE()` methods use a `WriteLock()` to make modifications to the list located at a bucket and `CONTAINS()` uses the corresponding `ReadLock` (see Section 8.3 in the text for a review of the `ReadWriteLock` concept). A `RESIZE()` method merely grabs all `WriteLocks` in sequential order (to avoid deadlock with a rival `RESIZE()` attempt) to halt activity during the course of the `RESIZE()` operation.
- b. **Lock-Free Closed-Address (LOCK-FREE)**: This design uses a regular `ReentrantLock` and does not require a `ReadWriteLock`, since the `CONTAINS()` method does not ever make use of the `ReadLock()` functionality. This design has the additional requirement that if a `CONTAINS()` and an `ADD()` or a `REMOVE()` method proceed concurrently that the result of the `CONTAINS()` is linearizable with the other call.
- c. **Linearly Probed Open-Address (LINEAR-PROBE)**: This is a standard linear probing style hash table, where each entry in the table has a counter which defines the maximum number of steps necessary to find a previously added item - that is, if an `ADD()` method had to walk  $k$  steps to find a vacant spot, the counter would be set to the max of the previous value and  $k$ . This allows both the `REMOVE()` and `CONTAINS()` methods to limit their searches. Beware of the potential for

deadlock when grabbing locks for multiple hash indices (as in `ADD()` and `REMOVE()`) in the event that the home index is not vacant. (For more, check out the Wikipedia article on [linear probing](#))

- d. **Cuckoo Hashing Open-Address** (CUCKOO-HASH): This design incorporates two hash functions - and two corresponding arrays to store items - to handle collisions. An item  $x$  can either be stored at location `table[0][ $h_0(x)$ ]` or `table[1][ $h_1(x)$ ]`, and the `ADD()` method swaps conflicting items between the two arrays in a chain in order to free up a slot for the given item, or resizes the table if the chain is too long or the table is too full to satisfy the request. (Section 13.4 of the text describes concurrent algorithms)
- e. **Your Design!** (AWESOME): This is your opportunity to design a really fancy version of either LOCK-FREE or LINEAR PROBE. In particular, you should extend the design in at least two non-trivial ways, demonstrating through experimentation of your own design why it is a good choice and providing your reasoning for its correctness. *Please include a section in your report where you concretely discuss the performance tradeoffs you made, your experiments to test those tradeoffs and your reasoning about the correctness of your algorithm.* You are free to pursue whatever design enhancements you like provided they are roughly as sophisticated as the following examples:
- **Concurrent Resize:** Design your hash table such that all operations, including `ADD()` and `REMOVE()`, proceed concurrently with `RESIZE()` method calls.
  - **Distributed Resize:** When a single thread performs the resize operation, there is the potential that if the rate of all threads making `ADD()` calls exceeds the rate of a single thread copying items, that the maximum speedup is governed by the single thread performing the `RESIZE()` operation. Instead, your design could offload the work of copying items to the other threads. Use your engineering judgment to decide among the myriad ways there are to do this - be sure to include your justification for both performance and correctness in the report.
  - **Lock-Free:** Can you build a concurrent hash table using only synchronization primitives like `GETANDSET()` and `FETCHANDINCREMENT` with no locks?
  - **Perfect Hashing:** One disadvantage of using a list for each bucket in a Closed-Address hash table is that the traversal of that list is inherently serial and takes time linear in the length of the list in the worst case. In the theoretical analysis of hash tables, there is a technique called **Perfect Hashing** which reduces the access time of a hash table to  $O(1)$  (see slide 16 [here](#) and page 4 [here](#) for more details). Essentially, the technique amounts to another hash table (albeit much smaller) to hold items at each bucket, rather than a list: a hash table of hash tables.
- f. **Your Application Specific Design** (APP-SPECIFIC): Suppose you knew the precise mix of `ADD()`, `REMOVE()` and `CONTAINS()` calls? What would you change in your design? Make a special-purpose design for the case where there are 20% `ADD()` calls, 20% `REMOVE()` calls and 60% `CONTAINS()` calls (each of which returns *true* 50% of the time) and the table is initialized with 1 million items. Remember, while the average number of entries may stay at 1 million, it could vary tremendously depending on the mercurial nature of our random number generator. So, while it may be tempting to make a table that is not resizable, you are not permitted to drop `ADD()` calls, so your code must handle this overflow condition gracefully.

## Assignment 2. Tests

We will test the concurrency of your hash table designs with a simple throughput test, taking packets from a `HASHPACKETGENERATOR` object which are encoded as `ADD`, `REMOVE` or `CONTAINS` requests. Then each packet is applied to the hash table according to this encoding, calculating the `FINGERPRINT()` - a proxy for generic packet processing work - as before. We will build three versions of this code:

- **SERIAL:** A single thread serially retrieves packets and makes the appropriate calls into a serial hash table implementation. This version of the code is provided and is configurable by the following parameters:
  - NUMMILLISECONDS ( $M$ ) the time in milliseconds that the experiment should run.
  - FRACTIONADD ( $P_+$ ) the percentage of packets that result in an ADD() call
  - FRACTIONREMOVE ( $P_-$ ) the fraction of packets that result in a REMOVE() call. Note that the fraction of CONTAINS() packets is  $1 - P_+ - P_-$ .
  - HITRATE ( $\rho$ ) the fraction of CONTAINS() calls that find data in the hash table.
  - MAXBUCKETSIZE ( $B$ ) the maximum allowed bucket size. If any bucket becomes this big, we have to resize the hash table.
  - MEAN ( $W$ ) the expected amount of work per packet.
  - INITSIZE ( $s$ ) the number of items to be preloaded into the table before the Dispatcher and Workers begin working (call GETADDPACKET() in the HASHPACKETGENERATOR class  $s$  times).
- **PARALLEL:** As in *pset4*, a Dispatcher thread will solely retrieve packets from the HASHPACKETGENERATOR and write them into a bank of Lamport queues (QUEUEDEPTH = 8). Using your best load balancing strategy from the last assignment, the  $n$  Worker threads will retrieve the packets from the queue bank and then make concurrent calls into the hash table. This version of the code is configurable by the same parameters as SERIAL plus:
  - NUMWORKERS ( $n$ ) the number of Worker threads.
  - TABLETYPE ( $H$ ) the name of a hash table type (eg. LOCKING, CUCKOO-HASH etc.) or NONE.

When  $H = \text{NONE}$ , the Workers should just drop the packets without processing them. The purpose of this "PARALLELNOLOAD" setting is to determine the maximum throughput of the Dispatcher.

The starter code will provide some basic components (e.g. utilities for timing code, a random packet generator etc.) and an example template of the top-level MAIN functions for each version of the code. In particular, the following classes are provided (in addition to the classes provided in *pset4*):

- SERIAL as described above
- HASHPACKETGENERATOR the random generator, which is instantiated with PERCENTADD, PERCENTREMOVE and HITRATE as described above.
- SERIALHASHTABLE a basic serial Closed-Address hash table implementation.
- SERIALLIST a basic serial linked list.

You are free to modify or rewrite any of these classes except for HASHPACKETGENERATOR.

### Assignment 3. Hash Table Throughput Experiments

Next, you will perform the following set of experiments across various cross-products of these parameters. Each data point is a measurement taken on a dynamic system (a computer...) and is thus subject to noise. As a result, some care should be taken to extract representative data - we would propose running some reasonable number of experiments for each data point and selecting the median value as the representative. Please use your own engineering judgment to decide how many trials you require. Setting the experiment time  $M$  to 2000 (ie. 2 seconds) should suffice to warm up the JVM and the caches for all of the following experiments. In each of the following experiments, we describe an experiment that you should perform and analyze. Please provide your hypotheses for the trends you see, supported by additional experiments of your own design if your hypotheses are not sufficiently substantiated by these experiments.

- a. **Dispatcher Rate** Run PARALLELNOLOAD with  $n = 8$  and  $W = 1$  to estimate the rate at which the Dispatcher retrieves packets and writes them into the queue bank. How does this rate compare with the same rate in *pset4* programming assignment 13d?
- b. **Speedup** Run SERIAL and PARALLEL with  $n \in \{1, 2, 4, 8\}$ ,  $W = 4000$ ,  $s = 0$ , all  $H$  except APP-SPECIFIC and under two load configurations:
  - (a) **Mostly Reads**  $P_+ = 0.09$ ,  $P_- = 0.01$ ,  $\rho \in \{0.5, 0.75, 0.9\}$
  - (b) **Heavy Writes**  $P_+ = 0.45$ ,  $P_- = 0.05$ ,  $\rho \in \{0.5, 0.75, 0.9\}$

Make 6 plots (one for each load configuration) with speedup (relative to SERIAL) on the  $Y$ -axis and  $n$  on the  $X$ -axis. Analyze this data and present your most insightful observations in the report, including the overhead of the parallel hash tables with  $n = 1$ .

- c. **Application Specific Test** Run SERIAL and PARALLEL with  $s = 10^6$  (the initial occupancy of the table),  $n \in \{1, 2, 4, 8, 16, 32\}$ ,  $W = 4000$  and all  $H$  under the following load:  $P_+ = 0.20$ ,  $P_- = 0.20$  and  $\rho = 0.5$ .

Try some different values by running on your own machine for the MAXBUCKETSIZE parameter in each test to see how it affects performance (there is a tradeoff between having frequent resizes and having larger chains). You can pick a value based on what you have observed, describe your intuition on why it seems right and use it for all the experiments on the class machine.

As usual, you should develop your code for correctness on your own machine, such that you can run these experiments to your satisfaction, prior to submitting the jobs using `cqsub`. **Do not write scripts to automate running multiple `cqsub` commands.** Instead, you may use `cqsub` to run a script which runs multiple `java` commands. Remember that you are timesharing with other students on the `eeecs-ath-34` server and taking turns running jobs through `cqsub`, so keep your jobs short but meaningful.

**Running the tests** Make sure the `.java` files are listed in the `makefile`. You can compile all classes by typing `make`.

We have written a script that you can use with `cqsub` to run your experiments in batch. It will save the experiment data in a Python pickled format, and also generate a `LATEX`code for the speedup plots. You will have to provide the maximum bucket size to the script. Also you should enter the packet rate/ms in the serial executions for the script to calculate the speedup. These variables are hard coded in the script. Make sure to read all the comments (and code) in the script before asking questions. If you significantly improve the code, we encourage you to share it with others.

The code is in Github: <https://github.mit.edu/hayks/6816-ps5-batch>

Note that the script will not explain why the plots look the way they do, you should add that bit!

#### Assignment 4. Report

Write a short but self-contained report summarizing your results from the experiments and the conclusions you draw from them, given your code and your answers to the questions above and including any additional experimental data you used to draw conclusions. That is, somebody should be able to read the report alone and understand what code you developed, what experiments you ran and how the data supports the conclusions you draw.

In addition, for each of the six hash tables you implemented, provide the reasoning (though not a formal proof) as to why your designs are correct. In particular, for each method in each design you should specify whether the method is wait-free, lock-free, deadlock-free or starvation-free and give the reasoning for it.

**Bonus Points: Competition** We are going to hold a performance competition of the "Your Design" AWESOME hash tables. After all the submissions are in, the TAs will test each student's AWESOME hash table on a secret set of parameters (run using `cqsub`). You will receive extra credit if your hash table is in the top fraction in terms of scalability, and (separately) receive extra credit if your hash table is in the top fraction in terms of absolute parallel speed.